

# Parallel Computing with OpenMP

Assignment Project Exam Help

<https://eduassistpro.github.io/>

Dr Paul Ric

<http://paulrichmond.shef.ac.uk> Add WeChat [edu\\_assist\\_pro](https://eduassistpro.github.io/) COM4521/



The  
University  
Of  
Sheffield.

 NVIDIA

GPU  
RESEARCH  
CENTER

# Last Lecture

- ❑ We looked at how to make programs fast on a single core
- ❑ **But** we didn't consider parallelism
- ❑ Guess what we are going today?

Assignment Project Exam Help

<https://eduassistpro.github.io/>

Add WeChat edu\_assist\_pro

❑ Multicore systems and OpenMP

❑ Parallelising Loops

❑ Critical Sections and Synchronisation

❑ Scoping

❑ Data vs Task Parallelis <https://eduassistpro.github.io/>

Assignment Project Exam Help

Add WeChat edu\_assist\_pro

# Multicore systems

❑ Multi-core CPUs are a shared memory system

❑ Each CPU has access to main memory

❑ Each CPU can access all of the memory (hence shared)

❑ Each CPU cores have their own L2 and L3 cache


❑ This can cause a lack of

❑ If one core modifies its data, it is not synchronised on other cores

Assignment Project Exam Help

<https://eduassistpro.github.io/>

Add WeChat edu\_assist\_pro




# OpenMP

## ❑ Open Multi-Processing Standard

- ❑ An API that supports shared memory programming in C, C++ and FORTRAN
- ❑ Cross platform support using native threading
  - ❑ Higher level than OS models and portable
- ❑ Is not suitable for distributed memory (MPI)

## ❑ It is not an automatic parallelization language

- ❑ Parallelism is explicitly defined and controlled by the programmer
- ❑ Requires compiler directives, a runtime library and environment variables



# OpenMP Compiler Directives

## ❑ Use of #pragmas

- ❑ If not understood by the compiler then they are ignored
- ❑ Does not require serial code to be changed
- ❑ Allows behaviour to be specified which are not part of the C specification

Assignment Project Exam Help

```
#include <stdio.h>
#include <omp.h>
```

```
int main()
{
    #pragma omp parallel
    {
        printf("Hello World\n");
    }
    return 0;
}
```

<https://eduassistpro.github.io/>

Add WeChat edu\_assist\_pro

# Extending OpenMP Hello World

```
#include <stdio.h>
#include <omp.h>

int main()
{
    #pragma omp parallel
    {
        int thread = omp_get_thread_num();
        int max_threads
        printf("Hello World (Thread %d of %d)\n", thread, max_threads);
    }
    return 0;
}
```

Assignment Project Exam Help

<https://eduassistpro.github.io/>

Add WeChat edu\_assist\_pro

```
Hello World (Thread 5 of 8)
Hello World (Thread 6 of 8)
Hello World (Thread 2 of 8)
Hello World (Thread 7 of 8)
Hello World (Thread 1 of 8)
Hello World (Thread 0 of 8)
Hello World (Thread 3 of 8)
Hello World (Thread 4 of 8)
```

# Fork and Join

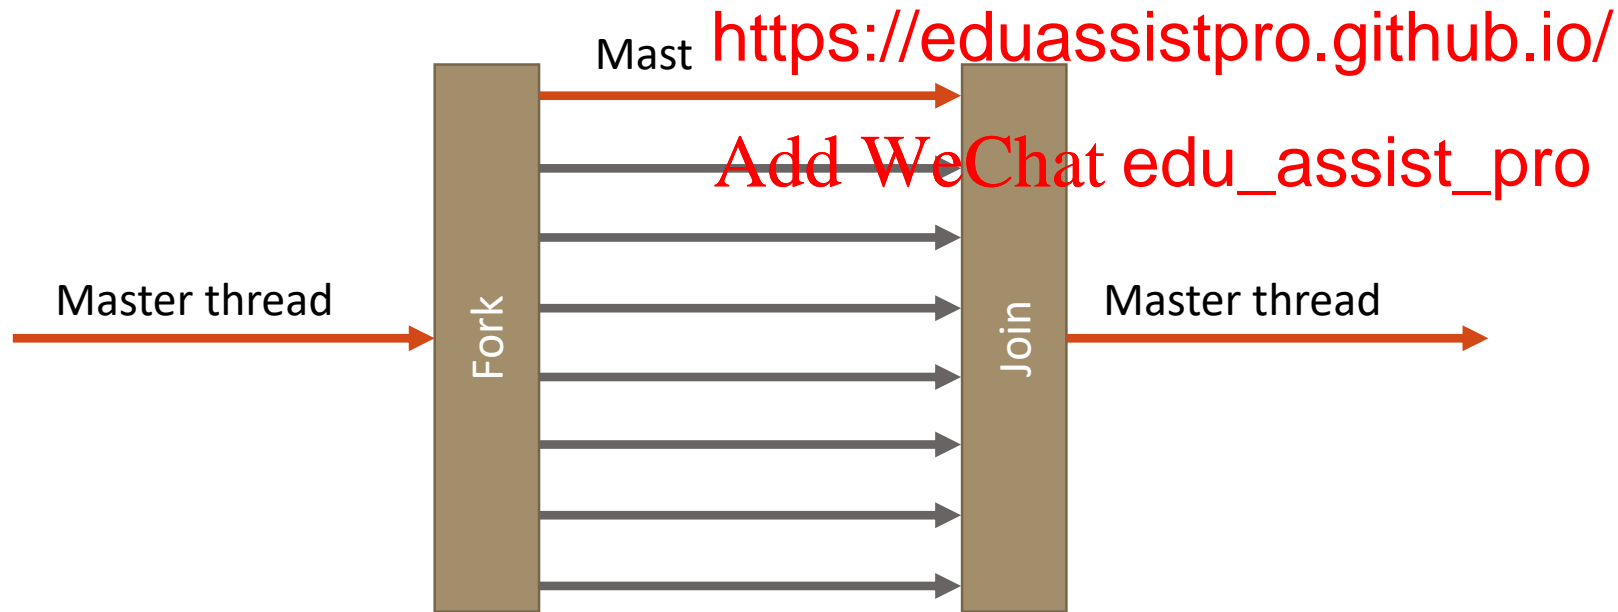
❑ OpenMP uses a fork a join model

❑ Fork: Creates a number of parallel threads from a master thread

❑ Master thread is always thread 0

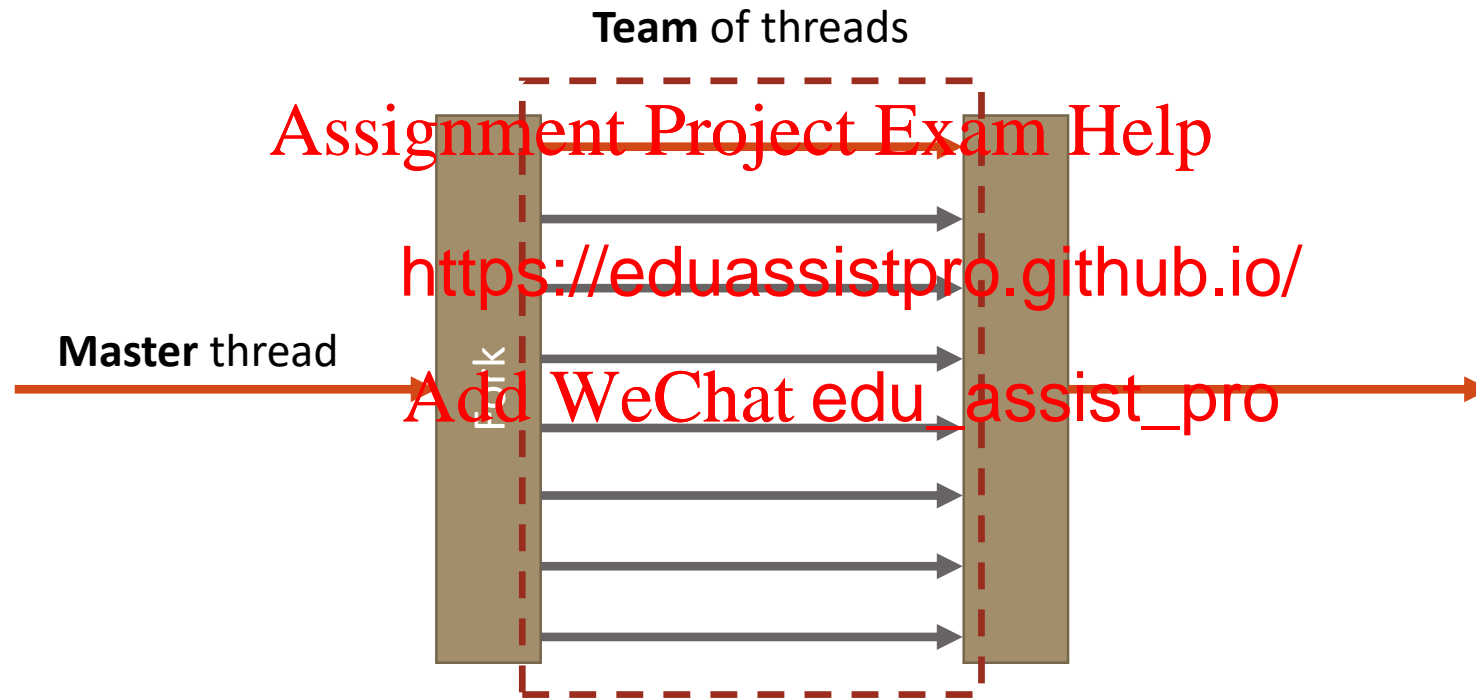
❑ No guarantee of order.

❑ Join: Synchronises thread termination and returns program control to master

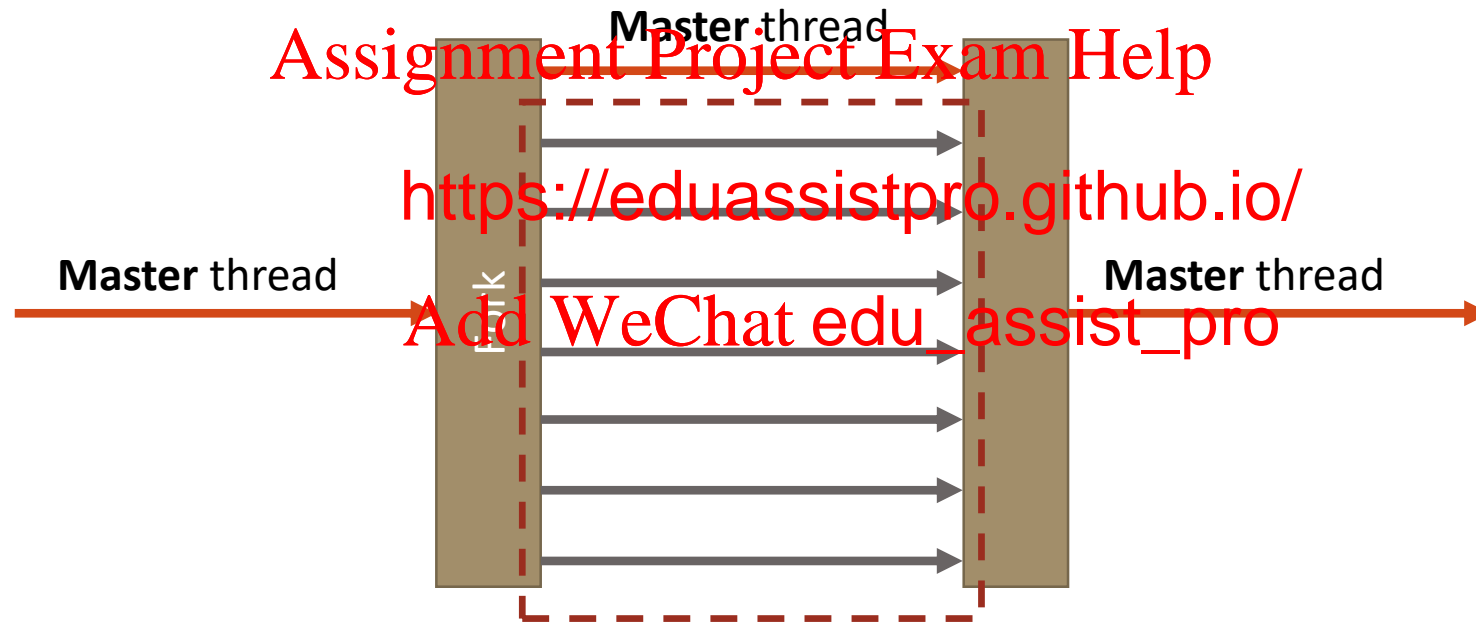




# Terminology



# Terminology



❑ Multicore systems and OpenMP

❑ Parallelising Loops

❑ Critical Sections and Synchronisation

❑ Scoping

❑ Data vs Task Parallelis <https://eduassistpro.github.io/>

Add WeChat edu\_assist\_pro

# OpenMP Syntax

## ❑ Parallel region directive

- ❑ `#pragma omp parallel [clause list] {structured block}`

- ❑ Spawns a number of parallel threads

## ❑ Clauses

- ❑ Are used to specify m directive e.g.

- ❑ Control scoping of vari

- ❑ Dictate the number of parallel threads (e

- ❑ Conditional parallelism

Assignment Project Exam Help

<https://eduassistpro.github.io/>

Add WeChat edu\_assist\_pro

```
#pragma omp parallel num_threads(16)
{
    int thread = omp_get_thread_num();
    int max_threads = omp_get_max_threads();
    printf("Hello World (Thread %d of %d)\n", thread, max_threads);
}
```

# num\_threads()

❑ Without this clause `OMP_NUM_THREADS` will be used

❑ This is an environment variable

❑ Set to the number of cores (or hyperthreads) on your machine

❑ This can be set globally by `omp_set_num_threads(int)`

❑ Value can be queried `threads()` ;

❑ `num_threads` takes <https://eduassistpro.github.io/> environment variable

❑ `num_threads()` does not guarantee the number requested will be created

❑ System limitations may prevent this

❑ However: It almost always will

Application

Compiler

Environment

OpenMP Runtime

Platform threading model  
(e.g. Windows threading or pthreads)



# parallel for

- ❑ #pragma omp for
  - ❑ Assigns work units to the team
  - ❑ Divides loop iterations between threads
- ❑ For can be combined with #pragma omp parallel for
  - ❑ Threads are spawned

```
int n;  
#pragma omp parallel for  
for (n = 0; n < 8; n++){  
    int thread = omp_get_thread_num();  
    printf("Parallel thread %d \n", thread);  
}
```

```
#pragma omp parallel  
{  
    int n;  
    for (n = 0; n < 8; n++){  
        int thread = omp_get_thread_num();  
        printf("Parallel thread %d \n", thread);  
    }  
}
```

<https://eduassistpro.github.io/>

Add WeChat [edu\\_assist\\_pro](#)

```
for (n = 0; n < 8; n++){  
    int thread = omp_get_thread_num();  
    printf("Parallel thread %d \n", thread);  
}
```

Which is the odd one out?

# parallel for

- ❑ `#pragma omp for`
  - ❑ Assigns work units to the team
  - ❑ Divides loop iterations between threads
- ❑ For can be combined with `assignment`
  - ❑ Threads are spawned

# Assignment for Project Parallel Exam Help

<https://eduassistpro.github.io/>

# Add WeChat edu\_assist\_5\_pro

```
#pragma omp parallel
{
    int n;
    for (n = 0; n < 8; n++){
        int thread = omp_get_thread_num();
        printf("Parallel thread %d \n", thread);
    }
}
```

```
Parallel thread 0  
Parallel thread 0  
Parallel thread 0  
Parallel thread 0  
Parallel thread 0  
Parallel thread 0  
Parallel thread 0  
Parallel thread 0  
Parallel thread 2  
Parallel thread 2  
Parallel thread 2  
Parallel thread 2  
Parallel thread 2  
read 2  
read 2  
read 2  
d 5  
d 5  
d 5  
Parallel thread 5  
Parallel thread 4  
Parallel thread 4  
Parallel thread 3  
Parallel thread 3  
Parallel thread 1  
...  

```

❑ Multicore systems and OpenMP

❑ Parallelising Loops

❑ Critical Sections and Synchronisation

❑ Scoping

❑ Data vs Task Parallelis <https://eduassistpro.github.io/>

Assignment Project Exam Help

Add WeChat edu\_assist\_pro





# What is wrong with this code?

❑ Consider a problem such as Taylor series expansion for  $\cos$  function

$$\cos(x) = \sum_{n=0}^{\infty} (-1)^n \frac{x^{2n}}{(2n)!}$$

$$\cos(x) = 1 - \frac{x^2}{2!} + \frac{x^4}{4!} - \frac{x^6}{6!} + \dots$$

Assignment Project Exam Help

```
int n;  
double result = 0.0;  
double x = 1.0;  
  
#pragma omp parallel for  
for (n = 0; n < EXPANSION_STEPS; n++) {  
    double r = pow(-1, n) * pow(x, 2 * n) / fac(2 * n);  
    result += r;  
}  
  
printf("Approximation is %f, value is %f\n", result, cos(x));
```

<https://eduassistpro.github.io/>

Add WeChat edu\_assist\_pro

# Critical sections

❑ Consider a problem such as Taylor series expansion for  $\cos$  function

$$\cos(x) = \sum_{n=0}^{\infty} (-1)^n \frac{x^{2n}}{(2n)!}$$

$$\cos(x) = 1 - \frac{x^2}{2!} + \frac{x^4}{4!} - \frac{x^6}{6!} + \dots$$

Assignment Project Exam Help

```
int n;  
double result = 0.0;  
double x = 1.0;  
  
#pragma omp parallel for  
for (n = 0; n < EXPANSION_STEPS; n++) {  
    double r = pow(-1, n) * pow(x, 2 * n) / fac(2 * n);  
    result += r;  
}  
  
printf("Approximation is %f, value is %f\n", result, cos(x));
```

<https://eduassistpro.github.io/>

Add WeChat edu\_assist\_pro

Multiple threads try to write to the same value!  
(undefined behaviour and unpredictable results)

# Critical sections

❑ Consider a problem such as Taylor series expansion for  $\cos$  function

$$\cos(x) = \sum_{n=0}^{\infty} (-1)^n \frac{x^{2n}}{(2n)!}$$

$$\cos(x) = 1 - \frac{x^2}{2!} + \frac{x^4}{4!} - \frac{x^6}{6!} + \dots$$

Assignment Project Exam Help

```
int n;  
double result = 0.0;  
double x = 1.0;  
  
#pragma omp parallel for  
for (n = 0; n < EXPANSION_STEPS; n++) {  
    double r = pow(-1, n) * pow(x, 2 * n) / fac(2 * n);  
    #pragma omp critical  
    {  
        result += r;  
    }  
}  
  
printf("Approximation is %f, value is %f\n", result, cos(x));
```

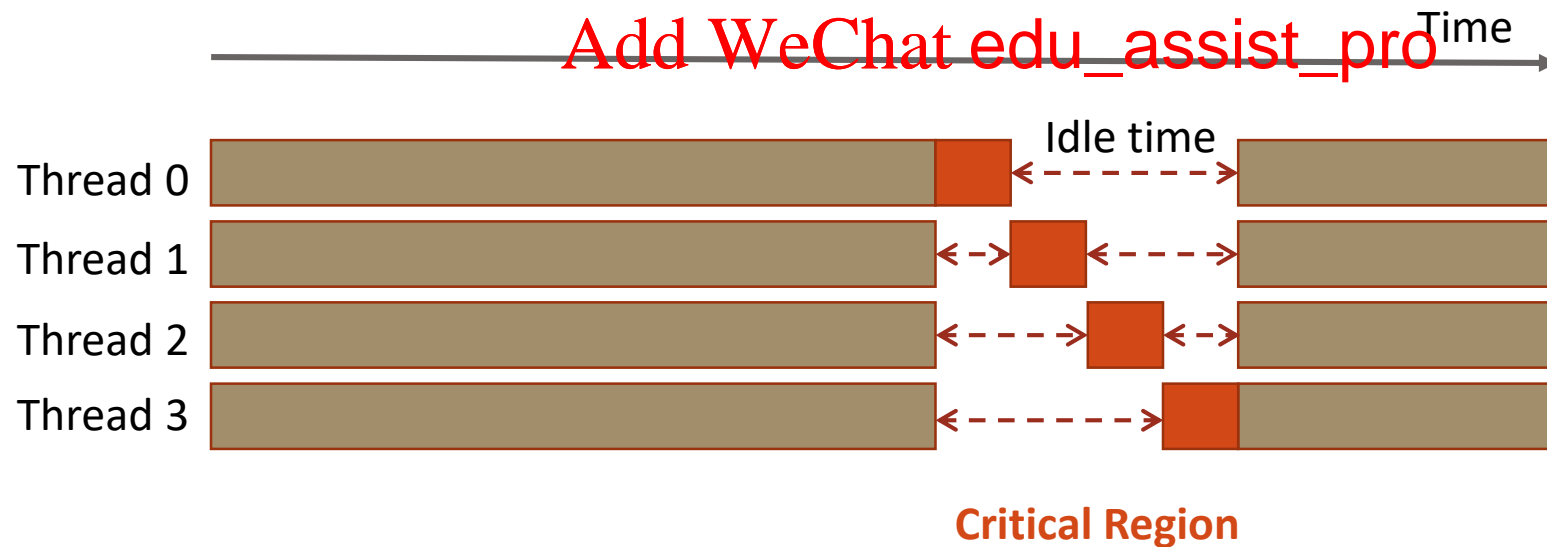
<https://eduassistpro.github.io/>

Add WeChat edu\_assist\_pro

Define as a critical section

# Critical sections

- ❑ `#pragma omp critical [name]`
  - ❑ Ensures mutual exclusions when accessing a shared value
  - ❑ Prevents race conditions
  - ❑ A thread will wait until no other thread is executing a critical region (with the same name) before being able to enter
  - ❑ Unnamed critical region: `#pragma omp critical`



# Atomics

- ❑ Atomic operations can be used to safely increment a shared numeric value
  - ❑ For example summation
  - ❑ Atomics only apply to the immediate assignment
- ❑ Atomics are usually faster than critical sections
  - ❑ Critical sections can be applied to general blocks of code (atomics can not)
- ❑ Example
  - ❑ Compute histogram of random values
  - ❑ Random is an `int` array of size `NUM_VALUES` containing random values within `0 : RANGE`
  - ❑ Histogram is an `int` array of size `RANGE`

Assignment Project Exam Help

<https://eduassistpro.github.io/>

Add WeChat [edu\\_assist\\_pro](#)

```
#pragma omp parallel
{
    int i;
    #pragma omp for
    for (i = 0; i < NUM_VALUES; i++){
        int value = randoms[i];
        #pragma omp atomic
        histogram[value]++;
    }
}
```

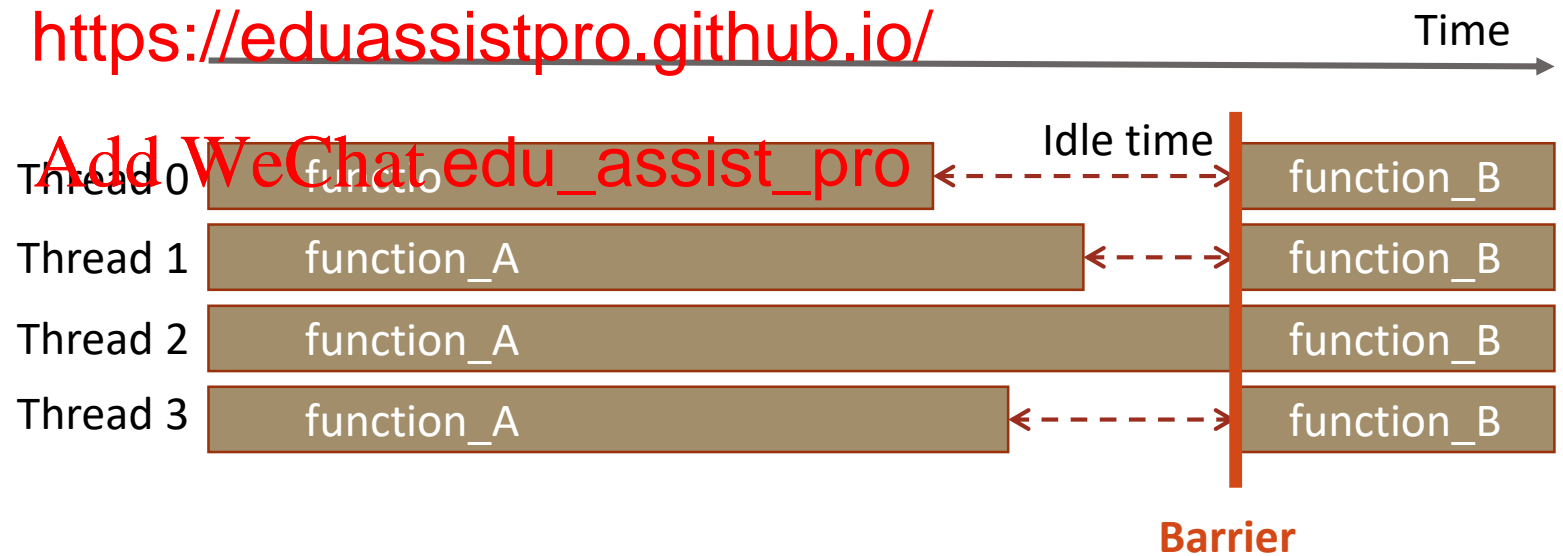
# Barriers

- ❑ `#pragma omp barrier`
  - ❑ Synchronises threads at a barrier point
  - ❑ Parallel regions have an implicit barrier
  - ❑ Can be used to ensure execution of particular code is complete
    - ❑ E.g. data read by func

Assignment Project Exam Help

<https://eduassistpro.github.io/>

```
#pragma omp parallel
{
    function_A()
    #pragma omp barrier
    function_B();
}
```



# Single and Master Sections

❑ `#pragma omp single { ... }`

❑ Used to ensure that only a single thread executes a region of a structured block

❑ Useful for I/O and initialisation

❑ First available thread will execute the defined region

❑ No control over which thread is used

❑ Will cause an implicit barrier

❑ E.g. `#pragma omp single`

❑ `nowait` will remove an implied barrier and is used to parallel for loops

❑ `#pragma omp master { ... }`

❑ Similar to `single` but will always use the master thread

❑ Is equivalent to using an `IF` clause

❑ E.g. `#pragma omp parallel IF(omp_get_thread_num() == 0) nowait`

❑ The `IF` clause makes the spawning of parallel threads conditional

❑ Preferable to `single`

❑ Does not require an implicit barrier

Assignment Project Exam Help

<https://eduassistpro.github.io/>

Add WeChat edu\_assist\_pro

# Master example

```
int t, r;
int local_histogram[THREADS][RANGE];

zero_histogram(local_histogram);

#pragma omp parallel num_threads(THREADS)
{
    int i;
    #pragma omp for
    for (i = 0; i < NUM_VALUES; i++) {
        int value = randoms[i];
        local_histogram[omp_get_thread_num()][value]++;
    }
    #pragma omp barrier
    #pragma omp master
    for (t = 0; t < THREADS; t++) {
        for (r = 0; r < RANGE; r++) {
            histogram[r] += local_histogram[t][r];
        }
    }
}
```

Same result as the  
atomic version

Assignment Project Exam Help

<https://eduassistpro.github.io/>

Add WeChat edu\_assist\_pro



❑ Multicore systems and OpenMP

❑ Parallelising Loops

❑ Critical Sections and Synchronisation

❑ Scoping

❑ Data vs Task Parallelis <https://eduassistpro.github.io/>

Assignment Project Exam Help

Add WeChat edu\_assist\_pro

# Scoping

- ❑ Scope refers to the part of the program in which a variable can be used
- ❑ OpenMP has different scoping to serial programming
  - ❑ We must define if a variable is private or shared between threads

<https://eduassistpro.github.io/>

- ❑ **Shared:** A variable can be accessed by all threads in the team
  - ❑ All variables outside of a parallel loop are shared by default
- ❑ **Private:** A Variable is local to a single thread and can only be accessed by this thread within the structured block it is defined
  - ❑ All variables inside a structured block are private by default

# Scoping

```
int t, r;
int local_histogram[THREADS][RANGE];

zero_histogram(local_histogram);

#pragma omp parallel num_threads(THREADS)
{
    int i;
    #pragma omp for
    for (i = 0; i < NUM_VALUES; i++)
        int value = randoms[i];
        local_histogram[omp_get_thread_num()][value]++;
}
#pragma omp barrier
#pragma omp master
for (t = 0; t < THREADS; t++) {
    for (r = 0; r < RANGE; r++) {
        histogram[r] += local_histogram[t][r];
    }
}
```

Shared

Assignment Project Exam Help But what about i?

<https://eduassistpro.github.io/>

Add WeChat edu\_assist\_pro

Private

# Scoping

```
int t, r;
int local_histogram[THREADS][RANGE];

zero_histogram(local_histogram);

#pragma omp parallel num_threads(THREADS)
{
    int i;
    #pragma omp for
    for (i = 0; i < NUM_VALUES; i++)
        int value = randoms[i];
        local_histogram[omp_get_thread_num()][value]++;
}
#pragma omp barrier
#pragma omp master
for (t = 0; t < THREADS; t++) {
    for (r = 0; r < RANGE; r++) {
        histogram[r] += local_histogram[t][r];
    }
}
```

Shared

i is private as it is  
the counter of the  
parallel for loop

Private

Assignment Project Exam Help

<https://eduassistpro.github.io/>

Add WeChat edu\_assist\_pro

# Explicit scoping

## ❑ Why is explicit scoping required?

❑ It is possible to use implicit scoping as in previous example

❑ Although it is good practice to use shared for any shared variables

❑ The clause default(shared or none) is helpful in ensuring you have defined variables scope correctly

❑ By changing the default variables and will give <https://eduassistpro.github.io/> enforces explicit scoping of

❑ const variables can not be explicitly (ays shared) - [more](#)

❑ Not enforced in windows but this is again

```
int a, b = 0;
#pragma omp parallel default(none) shared(b)
{
    b += a;
}
```

error C3052: 'a' : variable doesn't appear in a data-sharing clause under a default(none) clause

# Explicit scoping

## ❑ Why is explicit scoping required?

❑ Older C programming (C89) style has variable declarations before definitions and statements (including loops)

❑ Requires declarations to be made explicitly private for the parallel structured block

❑ E.g. Consider our atomic histogram example

Assignment Project Exam Help

<https://eduassistpro.github.io/>

Add WeChat edu\_assist\_pro

```
void calculate_histogram()
{
    int i;
    int value;
    #pragma omp parallel for private(value)
    for (i = 0; i < NUM_VALUES; i++) {
        value = randoms[i];
    }
    #pragma omp atomic
    histogram[value]++;
}
```

# Advanced private scoping

- ❑ If you want to pass the value of a variable outside of a parallel structured block then you must use the `firstprivate` clause
  - ❑ Private variables will be initialised with the value of the master thread before the parallel directive
- ❑ If you want to pass a private value to a variable outside of the parallel for loop you can use `use the loop`
  - ❑ This will assign the val

Assignment Project Exam Help

<https://eduassistpro.github.io/>

Add WeChat edu\_assist\_pro

```
int i = 10;
#pragma omp parallel private(i)
{
    printf("Thread %d: i = %d\n", omp_get_thread_num(), i);
}
```

```
Thread 0: i = 0
Thread 2: i = 0
Thread 1: i = 0
Thread 3: i = 0
```

```
int i = 10;
#pragma omp parallel firstprivate(i)
{
    printf("Thread %d: i = %d\n", omp_get_thread_num(), i);
}
```

```
Thread 0: i = 10
Thread 2: i = 10
Thread 1: i = 10
Thread 3: i = 10
```

❑ Multicore systems and OpenMP

❑ Parallelising Loops

❑ Critical Sections and Synchronisation

❑ Scoping

❑ Data vs Task Parallelis <https://eduassistpro.github.io/>

Assignment Project Exam Help

Add WeChat edu\_assist\_pro





# Data vs Task Parallelism

❑ Is an OpenMP parallel for clause data or task parallel?

Assignment Project Exam Help

<https://eduassistpro.github.io/>

Add WeChat edu\_assist\_pro

# Data vs Task Parallelism

❑ Parallelism over loops is **data parallelism**. i.e.

❑ The task is the same (the loop)

❑ Parallelism is over the data elements the loop refers to

❑ What about task parallelism?

❑ Task Parallelism: Divide and conquer

❑ This is supported by

❑ Further task parallelism is supported

❑ This is OpenMP 3.0 spec and not supported in Visual Studio 2017

❑ Very similar to sections

Assignment Project Exam Help

<https://eduassistpro.github.io/>

Add WeChat edu\_assist\_pro

# Sections

❑ `#pragma omp sections [clauses]`

❑ Defines a code region where individual sections can be assigned to individual threads

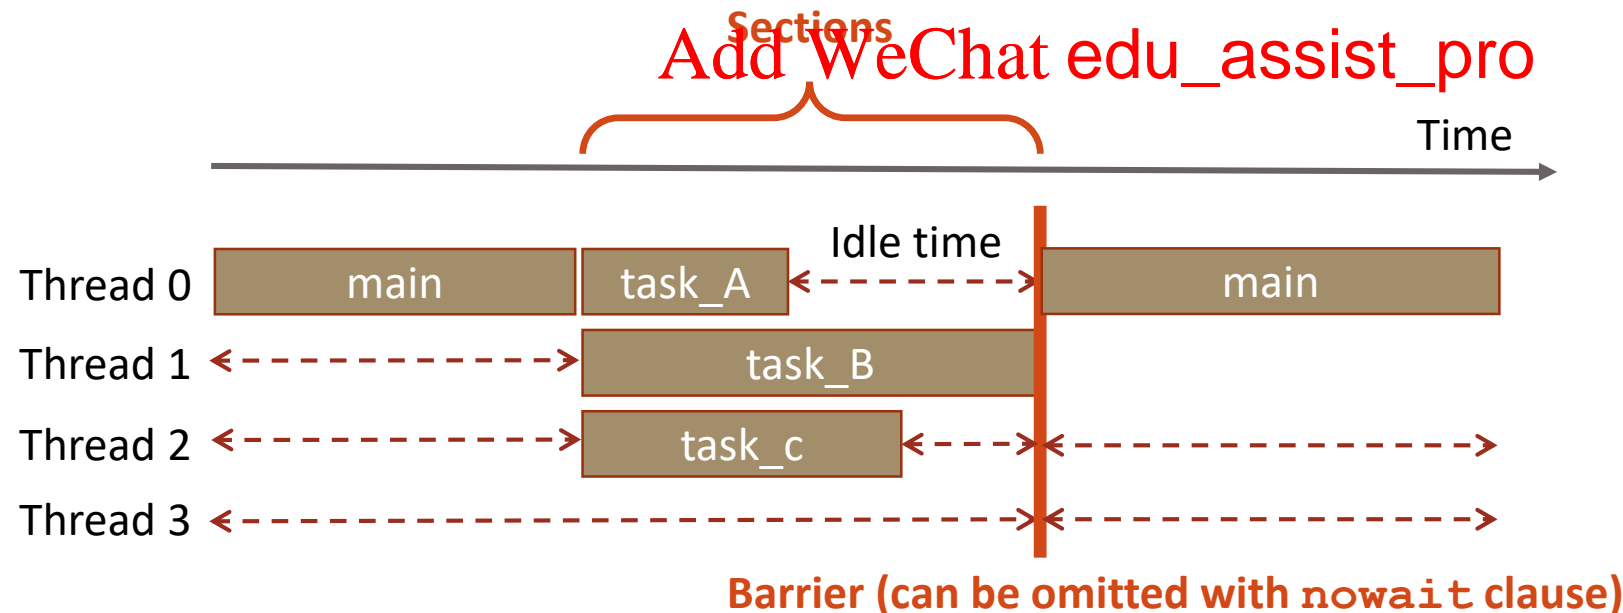
❑ Each section is executed exactly once by one thread

❑ Unused threads wait for

```
#pragma omp parallel
{
  #pragma omp sections
  {
    #pragma omp section
      task_A();
    #pragma omp section
      task_B();
    #pragma omp section
      task_C();
  }
}
```

Assignment Project Exam Help

<https://eduassistpro.github.io/>



# Sections

❑ If `nowait` clause is used then sections omit the barrier

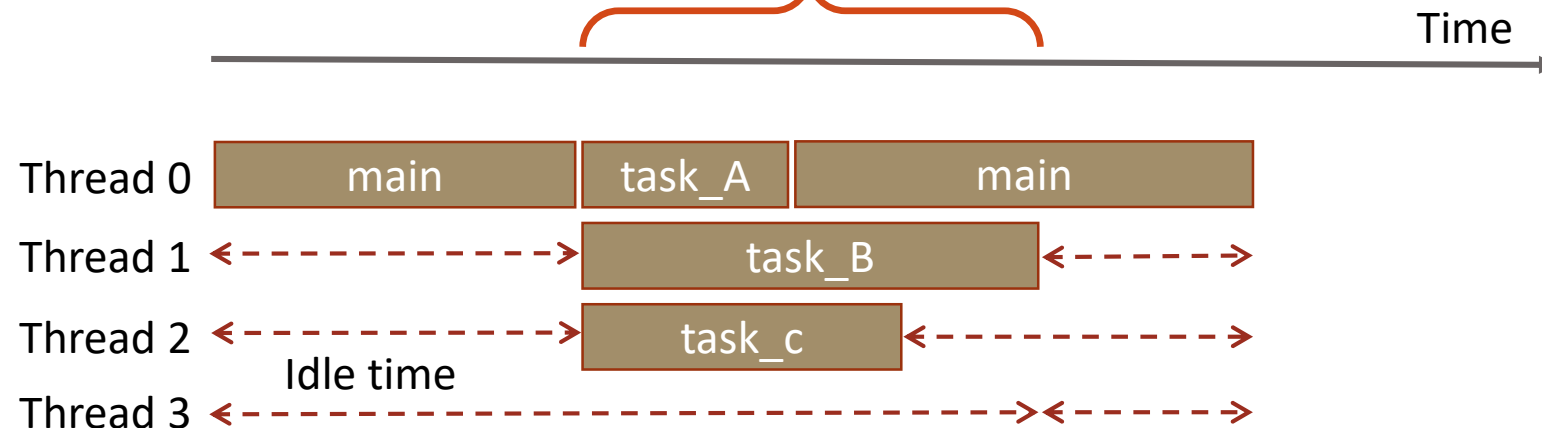
❑ will immediately enter other parallel sections

```
#pragma omp parallel
{
    #pragma omp sections nowait
    {
        #pragma omp section
        task_A();
        #pragma omp section
        task_B();
        #pragma omp section
        task_C();
    }
}
```

Assignment Project Exam Help

<https://eduassistpro.github.io/>

Sections  
Add WeChat edu\_assist\_pro



Barrier (can be omitted with `nowait` clause)

# Summary

- ❑ OpenMP lets us add explicit parallelism to serial code
  - ❑ We can parallelise loops or tasks
- ❑ OpenMP uses directives to modify the code
  - ❑ This enables to portability (serial and parallel code is the same)
- ❑ OpenMP exposes both serial and parallelism using a fork and join model
- ❑ Care must be taken on parallel blocks that require access to shared variables
  - ❑ There is a distinction between private and shared variables within a parallel blocks scope

Assignment Project Exam Help

<https://eduassistpro.github.io/>

Add WeChat edu\_assist\_pro